

Ogwi — Decision & Risk Log

Handover Document 4 of 4 · v2.1 (context edition) · August 2026 · Standalone. This log records approaches that were considered and **rejected** — they must not be built, however reasonable they look in isolation or however often the learning-science or gamification literature recommends them — plus deliberate deferrals and accepted risks. Documents 1–3 never depend on this log; it exists so future contributors don't reintroduce rejected ideas in good faith.

How to use this log, day to day: before proposing or building anything that resembles an entry here, read its reason first — most entries were genuinely good ideas that lost to a better principle, which is exactly what makes them dangerous later. If you believe circumstances have changed, the path is an owner decision, never a quiet build. The map: \$SA–D are product-mechanics rejections (motivation, learning, measurement, AI behaviour); \$E is deferred-with-trigger — build these only when the stated condition fires; \$F is accepted risk, held with open eyes; \$G–H are the commercial, social and brand perimeter.

A. Motivation & reward — do not build

- **Streaks, XP, points, coins, gems, badges-as-currency, leaderboards, login bonuses, daily-chest mechanics.** One motivation system exists (helium). Research on classroom gamification found badge/leaderboard/coin stacks *reducing* motivation over time (Hanus & Fox 2015 — the overjustification effect); the product's answer is a single system tied to genuine competence.
- **Selling or gifting rescue: streak freezes, rescue purchases, "watch an ad to save Oggi," system-granted litres of any kind.** Selling rescue mechanics is selling fear; litres are earned by learning only.
- **Wagering mechanics** (banking litres against an outcome, double-or-nothing). Losable-by-wager currency breaks "earned by learning only" and stages fear.
- **Per-user physics** (personal leak rates, boosts, exceptions) **and per-learner difficulty settings of any kind.** The moment two learners at 400ft mean different things, altitude stops being honest. Long-term progression rewards are cosmetic only.
- **Speed-pressure quizzes on fresh material; frenzy timers on new content.** Time pressure suppresses the effortful processing that makes retrieval work. (Learner-invoked timers and opt-in challenges on already-learned material remain legitimate.)
- **Countdowns, deadlines, decay warnings, "time remaining," red states, overdue counts.** Altitude is the only signal, ambient not anxious; pending work is date-phrased.
- **A "most flights" statistic, landing counts, or any surfaced metric that rewards touching down.** Comeback awards reward the return, never the fall.
- **Descent animation, narration, or a "your streak died" push.** Decline is discovered, never watched; touchdown is shown warmly in-app only.
- **Escrowing litres until topic completion.** Pending is a within-sitting display state; every exit pays out.

B. Learning mechanics — do not build

- **Pretesting / "have a go first" openers on unstudied material.** Despite literature support, sessions open with retrieval of *previously studied* material only.
- **Embedded questions or check-quizzes inside theory lessons.** Checking lives in the quiz lesson type; lessons use rhetorical questions and hypophora only.
- **In-the-moment retry on a wrong quiz answer.** Immediate re-guessing encourages answer-switching without processing; the explanation now, the end-of-quiz wrong-answer lap, and a variant days later replace it.
- **Mastery gates** — locking advancement behind a threshold, attempt counters, forced relearning loops. Mastery is measured continuously; nothing ever locks; no attempt counts exist anywhere.
- **A configurable personal target grade.** The qualification's real pass mark is the single bar for mastered states, opportunity ranking and readiness; a second target forces a distinction with no learner-meaningful difference.
- **Per-session goal-setting prompts, pre-reveal confidence ratings, post-session reflection check-ins.** The sole surviving self-assessment is the one-tap end-of-lesson confidence check, which feeds no scoring.
- **Diagnostics or prior-learning credit at onboarding.** Every learner starts the syllabus from the beginning; there is no placement test.
- **A mandatory recap/bridge formula at page ends; word counts that split indivisible worked examples.** Connected always, formulaic never; artifacts and indivisible examples sit outside the caps.
- **Interleaving unrelated subjects.** Mixing helps only across *confusable* material; sets mix only along human-confirmed confusable edges.
- **Mnemonic / "memory hook" generated blocks; per-learner interest-based rewriting; per-learner reading-level re-levelling; mind maps; generated audio lessons; narrated slide formats.** The lesson feed, fixed reading level with true blueprint terminology, career-relevance framing, and text-to-speech narration of the actual lesson are the approved forms.
- **Mutating Teaching Notes** (notes that update, reorder, tick or reveal progressively mid-session). The notes are a static, complete crib sheet; all progress signalling lives in the objectives panel.
- **A five-rung hint ladder.** Three rungs (hint → worked step → answer); longer ladders read as withholding.
- **Same-day pairing restrictions on learner-initiated practice.** The blur/Teach pairing rule binds the session composer only; learners may choose freely.
- **Housing blur/Teach inside the test builder.** A test is questions asked of you; activities are things you do; they live in a separate Activities area.

C. Assessment, measurement & data — do not build

- **Readiness as a score or grade prediction; a "ready by [date]" prediction.** Readiness is the odds of passing, with certainty; a ready-by date would stack pace, retention and calibration uncertainty into a promise the model cannot keep.

- **Revealing readiness on coverage alone, or a bare percentage with no stage context.** The three-item unlock checklist (topic + two-module breadth + one mini-mock) guarantees calibration from birth; numbers never appear naked.
- **Predictions computed from the softened display layer.** The kind display exists for motivation surfaces only; engines and readiness read live values.
- **Mass-rescheduling reviews when parameters or forecasts change.** Only future interval assignments use new settings; existing due dates stand (reschedule-on-change causes due-pile shocks).
- **Per-learner scheduler optimization at launch.** Deferred (see E).
- **Hours studied, videos completed, login counts, attempt counts — measured or displayed.** These flatter ineffective study and create counts of shame; last-active exists internally only.
- **Exam-date-based features and awards** (final-approach countdowns, exam-day badges, exam-date reminders) — and exam-booking nudges or targets of any kind: the readiness threshold ends in the single celebration only. The product never asks for, stores, or uses exam dates.
- **Read-appending event logs** (evaluation records; state that mutates on view). History must record what happened, never that somebody looked; the one exception is the atomic materialisation of a transition first discovered by a read.
- **A push-allowance reset on any study.** The allowance resets only when fill climbs back above 100ft; micro-top-ups below the line must not re-arm notifications.

D. AI behaviour — do not build

- **A distress-detection classifier with crisis-resource signposting, wellbeing check-ins, or any therapeutic register.** Oggi's boundary is two fixed deflections: platform matters → the support desk; personal, emotional, medical or legal territory → a fixed, warm "I don't have the capabilities for that," with no advice clause. He is never a counsellor and never pretends to be.
- **A greeting that opens with a question requiring an answer, or any blocking prompt at session start.** The greet is at most a one-line invitation with a visible alternative; input is free from the first moment.
- **Manufactured callbacks** ("remember when..." without real evidence). Memory is performed only where engine data makes it real.
- **Answer-extraction detection, accusation messaging, or an "extraction rate" metric.** Ladder discipline is the whole defence; the product never polices the learner.
- **Silent memory writes; storing sensitive personal disclosures; mining chat transcripts into future context.** Told-notes require an in-chat acknowledgement; a sensitive-topics filter blocks those writes structurally; transcripts are short-term, for reports and QA only.
- **Oggi ending a conversation, contacting anyone outside the app (beyond the single flight push), or alerting a Wingman to anything.** Chosen disclosure only.
- **Role-play scenario lessons; arcade "you vs Oggi" race modes.** Rejected on cost/benefit and on the speed-pressure principle respectively. (The rejection is the *lesson type*; Oggi may still improvise a scenario on request in ordinary chat.)

E. Deferred by design (build only on the stated trigger)

- **Per-learner scheduler parameter optimization.** Launch runs on FSRS's published defaults (the median fit across ~20,000 real learner collections, benchmarked on hundreds of millions of logged reviews; slightly *more* accurate for binary graders). Build only if cohort prediction error (RMSE) sustains above 0.07 for a month, or a segment shows systematic misprediction. Safe to defer because no-mass-reschedule makes later adoption migration-free.
- **The full severity-ranked QA queue.** Launch runs the lean loop (one inbox, weekly triage, urgent pause route). Build the full queue only on sustained report volume or anomaly rates above the configured thresholds for a month.
- **Multi-qualification journeys.** Deferred until the first published pass rate exists.

F. Accepted risks (open eyes, by decision)

- **No provenance or licence data is stored for imported question banks.** Rights clearance is wholly an out-of-band team responsibility; the system cannot later answer "where did this question come from." Mitigation: operational rights checklist before any import batch.
- **The Orbit display exception.** Above the 1000ft cap the altitude readout shows "??? ft" — the single, deliberate exception to the exact-figures principle, scoped to a celebration; the litre readout stays exact and all records carry the honest number. No further exceptions may cite this one as precedent.
- **The kind display layer.** Published mastery may transiently sit above live values (bounded by a 7-day half-life). Accepted because no prediction reads it; the invariant suite enforces the boundary.
- **Scheduler-efficiency claims are simulation-derived.** The "fewer reviews for equal retention" evidence behind the scheduling algorithm comes from large-scale simulation on logged data, not controlled live trials. Any marketing use must pass the source-check policy and state claims accurately.
- **Readiness calibration depends on self-reported exam outcomes.** Response bias is possible; treat the calibration KPI directionally until volume allows adjustment.
- **A single push threshold (~75L) tuned pre-launch.** If the weekend-heavy cohort grounds more than once a week *and* churns, the tuning order is: raise the proportional threshold toward 250ft or soften the proportional leak toward 45% — before touching anything else.

G. Additional product boundaries — do not build

Positioning, audience and commercial surface. Only the approved True North (*Make trusted career qualifications accessible and achievable for everyone*), the approved tagline (*"Going up."*) and the career-changer audience may be used; all alternate or historical True Norths, taglines and audience positionings are retired. "Duolingo for professional qualifications" / "RevisionDojo for professional qualifications" are internal orientation shorthand only — never external positioning lines. Do not build: a formal tagline-trademark-search workflow as a product task; competitor-comparison positioning; career-fit quizzes, taster experiences or "week in the life" previews (selection is an open menu); career-pathway

mapping; work-experience, CV, coaching or employer-matching services; B2B or team products in their entirety (per-seat licences, employer dashboards, compliance reporting); exam-fee or total-cost-of-qualification modules; book-a-call or any sales-assisted purchase motion; £25/month or any price point other than the \$39/month base; a day-count free trial or its "survive one missed day" rule (the trial ends at the first completed session); the borrowed-statistics ban and the single-training-statistic rule as standalone policies (both superseded by the universal source-check policy); the PWA-form install nudge (superseded by the native app-store nudge).

Flight and arcade variants. Do not build: banked-first mastery framing of altitude (altitude is current pace, never banked history); weekly race-day sessions; bedtime-silent Oggi variants; weather or storm declarations in the sky; a plain-mode Flight toggle; altitude plateaus or decay-freezing while the app is open; speed rounds with leaking-helium timers; arcade as a lesson type. **Frenzy/multiplier time-windows of any kind — including learner-invoked ones — are not built:** the approved drama layer is exactly the hard-question doubles, combo chains and quiz-scoped multipliers inside quizzes; learner-invoked *plain* timers and cosmetic seasonal challenges are the only time-based opt-ins.

Social, accountability and marketing concepts. Do not build: the predecessor accountability-link system in full (public links showing mastery, last-active or time-spent); a shareable "branch holds" readiness summary; relationship chips on Wingman surfaces; digest co-authoring; automatic downgrade to fortnightly digests; encouragement timing intelligence (encouragement always delivers instantly); the five-supporter "Ground Crew" shared-feed tier; Formation Flying; household rooms; public digests from Tracker pages; gating the Tracker behind a mastery milestone; proactive human outreach of any kind (gate-stall or plateau contact — the desk is purely reactive); physical trust-signal modules; current-student or graduate Q&A programmes; alumni-guide or near-peer ladders; post-pass CV-bridge content; the standalone "5 S" website programme as a workstream.

Interface and feature boundaries. Desktop Oggi expands as a side panel; a central popup is permitted only as a narrow-width fallback. Standalone non-timed practice feedback never displays elapsed time (time appears only on exam-format runs and learner-timed tests). "Try Again" is not a blurt feedback option — Correct Mistakes and Continue are the two follow-ups. No human handoff exists inside Oggi chat (the boundary is the fixed deflection to the support desk). No "stormy week ahead" or any weather declaration; no bedtime-silent audio default. No Oggi skins or restyles as cosmetic unlocks — Oggi keeps the brand look (cosmetics are city themes and award avatars).